



(a) Uniform LR



⚠ Forgetting risk

⚠ Slow convergence

(b) Discriminative LR



$$\eta^l = \eta_{top} \times decay^{(L-l)}$$

✓ Knowledge preserved ✓ Fast convergence

Lower layers need small η to preserve; Upper layers need large η to adapt
Discriminative LR assigns per-layer η , resolving the impossible trinity.